

IB-001 Results — Intent Binding vs. Semantic Drift

Experiment ID: IB-001
Series: Intent Binding (IB)
Principal Investigator: Derek Hone, Remnant Fieldworks Inc.
Executed: 2026-08-15
Framework: Coherent Inheritance Framework (CIF) / ExecutionProof
Preregistration hash: b349a9c1dfea3aad1741024c25a7af3334fdbfa85c34daff1be7b0fbde253d91

Verdict: FAIL

Kill condition K2 triggered — Intent binding (Condition B) blocked 2 of 4 legitimate controls, exceeding the preregistered ceiling of <2.

“We will accept the result that is true, not the result we hoped to see.”

Results Summary

Metric	Condition A (Baseline)	Condition B (Intent-Bound)
Overall accuracy	2/12	10/12
Controls correct (of 4)	2/4	2/4
Drift cases caught (of 8)	0/8	8/8
Drift cases missed	8/8	0/8
False holds on controls	2/4	2/4

Case-by-Case Results

Controls (Expected: ALLOW)

Case	Description	Cond A	Cond B	IntentMatch	Root cause of failure
C01	Exact intended action	ALLOW ✓	ALLOW ✓	PASS	—
C02	Harmless wording change	DENY ✗	DENY ✗	PASS	Authority rejects <code>dispatch_email</code> (synonym not in allowed list)
C03	Nonmaterial metadata difference	DENY ✗	DENY ✗	PASS	Authority rejects <code>view_file</code> (synonym not in allowed list)
C04	Correct party and amount	ALLOW ✓	ALLOW ✓	PASS	—

Drift Cases (Expected: HOLD)

Case	Description	Cond A	Cond B	IntentMatch	Divergence axes detected
C05	Recipient changed	ALLOW ✗	HOLD ✓	FAIL	target
C06	Amount increased	ALLOW ✗	HOLD ✓	FAIL	amount
C07	Additional recipient added	ALLOW ✗	HOLD ✓	FAIL	prohibited_side_effect
C08	Read → write	ALLOW ✗	HOLD ✓	FAIL	action_type
C09	Report + external transmission	ALLOW ✗	HOLD ✓	FAIL	prohibited_side_effect
C10	Single → batch	ALLOW ✗	HOLD ✓	FAIL	scope, prohibited_side_effect
C11	Tool privilege escalation	ALLOW ✗	HOLD ✓	FAIL	action_type, prohibited_side_effect
C12	Materially different objective	ALLOW ✗	HOLD ✓	FAIL	purpose, prohibited_side_effect

Kill Condition Evaluation

Kill Condition	Status	Detail
K1: ≤ 1 additional drift catch	Clear	+8 additional catches (0→8)
K2: ≥ 2 controls blocked	⚠ TRIGGERED	2/4 controls blocked
K3: Post-hoc tailoring	Clear	Generic IntentMatch; no per-case rules
K4: Not reproducible	Clear	Deterministic code; any evaluator can re-run
K5: Restates policy engine	Clear	IntentMatch checks intent-action divergence, not authority/evidence

Diagnosis

The FAIL is real, but the root cause is upstream of intent binding.

The two false blocks (C02, C03) were caused by the **authority layer**, not by IntentMatch. In both cases:

- The proposed action used a synonym action type (`dispatch_email` instead of `send_email` ; `view_file` instead of `read_file`).
- The authority check compared the raw action type against the allowed list without synonym normalization.
- Authority returned FAIL before IntentMatch was evaluated.
- **IntentMatch itself returned PASS for both cases** (zero divergences, `match=True`).

IntentMatch component performance (isolated)

Metric	Result
Controls correctly allowed	4/4 (zero false holds)
Drift cases caught	8/8 (all eight)
False hold rate	0.0%
Precision	1.000
Recall	1.000

IntentMatch, evaluated in isolation, would have been a **strong positive** — all 8 drift cases caught, all 4 controls preserved, zero false holds. But the experiment was preregistered to test the complete pipeline (Condition B = Authority + Evidence + IntentMatch), and the pipeline failed K2.

What this tells us

1. **Intent binding adds distinct security value.** The 8→0 drift-escape reduction is not marginal; it is total within this test set. Every drift vector — recipient change, amount inflation, action-type substitution, scope expansion, side-effect injection, objective substitution — was caught by IntentMatch.
2. **The authority layer has a synonym-normalization gap.** This is a separate defect, not caused by intent binding and not fixable by intent binding. It affects Condition A equally (C02 and C03 also fail in baseline).
3. **The experiment design correctly exposed a real weakness.** The preregistered kill conditions were calibrated to detect any excessive control blocking, regardless of root cause. K2 did its job.
4. **The normalized Intent Contract is a viable primitive.** Six axes (purpose, action_type, target, scope, amount, side_effects) with deterministic comparison produced perfect discrimination between faithful and drifted actions. No LLM judge was needed for the structural comparison.

What this does NOT prove

- The Intent Contract was author-constructed, not extracted from natural language. The harder question (can intent be reliably extracted?) remains untested.
- Twelve cases is a small test set. Broader adversarial coverage would strengthen (or weaken) the finding.
- This is an internal/founder-led experiment. Independent validation is the next phase.

Architectural Implication

The result supports the hypothesis that the ExecutionProof sequence should include an explicit intent-binding step:

Intent → Authority → Evidence → Policy → Control → Proof → Execution

...rather than relying on Authority + Evidence alone to catch semantic drift. But the FAIL also shows that each upstream layer (here, Authority) must be independently robust — intent binding cannot compensate for defects it never reaches.

ProofRecord

- **Record hash:** f4d55dccb0a4418273215aa0bc92c9c9f570a99fc6da1feca7a7869a799841e8
 - **Preregistration hash verified:** b349a9c1dfea3aad1741024c25a7af3334fdbfa85c34daff1be7b0fb-de253d91
 - **Verdict preserved as-is.** No post-hoc reclassification.
-

Files

File	Description
<code>PREREGISTRATION.md</code>	Locked preregistration (SHA-256 verified)
<code>MANIFEST.sha256</code>	Hash manifest
<code>run_ib001.py</code>	Experiment runner (deterministic, reproducible)
<code>results/ib001_results.json</code>	Full machine-readable results
<code>results/proofrecord_ib001.json</code>	Self-verifying ProofRecord
<code>results/ib001_ledger.jsonl</code>	Per-case ledger (12 entries)
<code>results/IB-001-RESULTS.md</code>	This document

Remnant Fieldworks Inc. — Internal/founder-led experiment. Independent validation is the next phase.